

CS250B: Modern Computer Systems

Lab 1: Matrix Multiplication Derby

Due: 2021-05-13

Overview

In this lab, you will write a best-effort multithreaded implementation of a dense matrix-matrix multiply in C/C++, and compare it against other implementations with various levels of optimizations.

Files Involved

In the lab directory, you will find four files, **guide.pdf**, **main.cpp**, **mult.cpp**, and **Makefile**. You will only edit **mult.cpp**. **guide.pdf** is this document

Implementation guide

Try to use all ideas discussed in class, including AVX, cache-obliviousness and loop unrolling. Some ideas may not work well together. Achieve the best you can!

The sole function in the **mult.cpp** file, **mult**, takes in five arguments: **a**, **b**, **c**, **matrix_size** and **thread_count**. **a** and **b** are to be multiplied and stored in **c**. **matrix_size** defines the size of the matrix input, where there is **matrix_size**² elements in the square matrix. **thread_count** define the number of threads to be spawned.

The **mult** function is currently implemented using a naive triple-loop method.

In order to build and run the project run **make**, and then **./matrix T M**. **T** is the number of threads argument that will be passed to your **mult** function. **N** is the size of the matrix, where the number of elements in the matrix is **N*****N**. The default value of **N** is 2048 unless specified.

Evaluation environment

Your code will be tested on various (but reasonable, desktop-class) machines, with different cache, processor and memory characteristics, using a different thread values. The matrix size will vary between 2048 to 16384.

All machines will have AVX2 and FMA support at least.

The derby

Each submission will be given a code name for anonymity, and I will create charts and leaderboards to compare various implementations, as well as some implementations I will make with various levels of optimizations.

Grading will not be on a curve, so relax, and do your best